Title

H. Sherry Zhang
Department of Statistics and Data Sciences, University of Texas at Austin
and
Roger D. Peng
Department of Statistics and Data Sciences, University of Texas at Austin

June 4, 2026

Abstract

Current reproducibility practices have greatly improved our ability to replicate published results, but they do not necessarily make data analyses easier to understand. Audience seeking to comprehend an analysis need to piece together information scattered across manuscripts and code scripts to understand the sequence of analytical steps and the rationale behind them, with some details may not explicitly recorded in either source. Exploring alternative analytical choices is even more challenging, as it requires understanding and modifying code that may be unfamiliar and consider how it would impact upstream and downstream code. To address these challenges, we propose a schema that represents data analyses as a collectino of analytical decisions, together with their justification, inputs, and outputs. Coupled with large language models (LLMs), the schema can be inferred from natural language descriptions and translated into analysis scripts. The structured representation also enables exploration of alternative analytical choices, allowing multiple analytical workflows to be generated and compared. We demonstrate the proposed workflow using the many-analysts football study, where LLMs create the schema from the decision chocies in team reports, generate corresponding analysis scripts, and explore alternative random effect specifications. As an intermediate representation between the manuscript and code scripts, our schema makes analyses more comprehensible and extendable from a reader's perspective.

Keywords: Reproducibility, Large Language Models (LLMs), Data analysis

1 Introduction

Data analysis provides crucial evidence to support scientific claims in publications. To ensure that results can be independently scrutized and replicated, reproducibility practices have long adcovated making data and code publicly available whenever possible. These practices have greatly improved the reproducibility of published results across scientific domains and have been required as part of journal submission processes.

An implicit assumption underlying current practices is that reproducibility implies understandability, given that all materials required to reproduce the results are available to readers. Specifically, understandable implies that an independent researcher should be able to, without much difficulty, track the steps that have been done and understand why each step was taken. In practice, however, as modern data analyses grown in complexity, with data often drawn from multiple sources and involving potentially spatial, temporal, and multivariate aspects, reconstructing an analysis is rarely straightfoward. Instead, it is often a frustrating process of trial and error, where there is little feedback when results do not match those reported. What's available for readers, the manuscripts and code scripts, are usually not perfect for understanding an analysis, but different communication goals. A manuscript is a curated narrative designed to present new scientific claims, rather than a detailed log of the data analysis workflow. Code scripts, on the other hand, show all the computational details, but don't inherently carry the underlying scientific rationale behind each step. Additionally, code scripts reflect author's choice of programming language, personal coding style, and customized abstractions built to streamline repetitive analytical tasks, specific to the project. These create additional difficulties for readers to understand the code and extend the analysis. Hence, there is a distinction between an analysis being machine-replicable and being understanble to readers. In practice, readers often spend

substantial time reconstructing analytical decisions by piecing together information from manuscripts and code scripts, even when all required materials are available. Recent advances in LLMs have shown strong capabilities in generating code from natural languages and repairing erroneous syntax, opening new opportunities for making data analysis easier to comprehensible, modifiable, and adaptable.

In this paper, we present a schema designed to describe data analysis in a structured way that captures the decisions and their justifications, as well as the inputs and outputs of each step. We also demonstrate how this schema, coupled with LLMs, can be used to generate code scripts from natural-language descriptions in manuscripts, and explore alternative decision choices. We apply our schema to the many-analyst football study, where LLMs create the schema from the decision choices reported in the final reports of each team, and generate and validate the corresponding code scripts against the original published results. We further demonstrate how to the schema supports specification and exploration of alternatives random effect specification for Team 5.

The paper is organized as follows: in Section 2, we review the relevant literature on reproducibility and communicaition of data analysis, as well as the application of LLMs in data science. In Section 3, we present our methodology, including the schema representation for describing data analysis and the workflow for generating and validating code from the schema by LLMs. In Section 4, we apply our methodology to the many-analyst football and demonstrate that we can reproduce the results reported in the paper, identify errors in the original reports, and highlight failure modes of current LLMs. Finally, in Section 5, we discuss the implications of our work and potential future directions.

2 Literature review

2.1 Reproducibility and communication of data analysis

The reproducibility movement has been a major development in modern scientific research, motivated in part by concerns regarding the replication of published findings and the increasing reliance on computational methods across disciplines. In statistics, advocates of reproducible research argue that scientific claims should be accompanied by the data and code necessary to regenerate reported results (Peng 2011). This movement has led to the adoption of platforms for hosting research artifacts, including GitHub, Figshare, Open Science Framework (OSF), and Zenodo. Literate programming tools such as Sweave (Leisch 2002), R Markdown (Xie 2015), Jupyter Notebooks (Kluyver et al. 2016), and Quarto (Allaire et al. 2026), have further promotes the integration of narrative and computation within a single document. These practices have increasingly been incorporated into journal policies through data and code availability requirements. For example, the *Journal of the American Statistical Association* introduced an associate editor for reproducibility and a reproducibility award to encourate and recognize these practices (Wrobel et al. 2024). Similar requirements have emerged across disciplines and have continued to improve the reproducibility of published research (Kohler et al. 2026, Xu & Yang 2026).

While these developments have improved the ability to reproduce computational results, they are not designed to make analyses easier to understand. Reproducibility concerns whether a result can be regenerated given the data, code, and computational environment, whereas comprehension concerns whether the audience can reconstruct the analytical decisions made by the authors, understand the rationale behind those decisions, and confidently modify the code to explore alternative choices. This distinction becomes increasingly im-

portant as data analyses grow in complexity. Decisions related to data preprocessing and exploratory data analysis are often reported only briefly in manuscripts and left as technical details for readers to infer from supplementary materials. Many decisions are not made in a linear fashion but emerge through an iterative process (Zhang & Peng 2025, Jun, Birchfield, De Moura, Heer & Just 2022, Jun, Seo, Heer & Just 2022), and it is not entirely clear how this iterative nature of data analysis can be communicated. Making changes to an existing script can often be difficult because analytical decisions are intertwined with implementation details, project-specific abstractions, and dependencies across different parts of the workflow, requiring readers to consider how modifications propagate through upstream and downstream components. As a result, readers seeking to understand, modify, or extend an analysis must often reconstruct the underlying rationale by piecing together information scattered across manuscripts, code scripts, and supplementary materials.

2.2 Large Language Models for Data Analysis

Advances in Large Language Models (LLMs) have led to widespread use of these systems for code generation tasks, including code completion and program repair (Jiang et al. 2026, Shah et al. 2026). Early work on code-oriented LLMs demonstrated strong capabilities in translating natural language descriptions into executable programs, with models such as Codex achieving substantial improvements on benchmark programming tasks (Chen et al. 2021, Lai et al. 2023, Huang et al. 2022). These capabilities have led to broad adoption of coding assistants thorugh web interfaces, APIs adn integrated development environments. In the R ecosystem, packages such as `ellmer` (Wickham et al. 2025), `chatLLM` (Owusu Boakye 2026*a*), and `LLMAgentR` (Owusu Boakye 2026*b*) provides increasingly accessible interfaces for access to latest LLM models.

Building on these advances, recent work explores fully automated systems that use multiple LLM agents to plan, generate, execute, and refine data analyses (Jansen et al. 2025, Hu et al. 2024). Our work relies on the same underlying capability of translating analytical specifications into code, but differs in its view of what should be automated. Data analysis is often underdetermined, with many plausible choices regarding data preprocessing, transformation, aggregation, and model specification (Gao & Xiao 2026, Bertran et al. 2026, Luo et al. 2025), by contrast, once these analytical decisions have been specified, translating them into executable code is often a comparatively straightforward task. This distinction motivates our focus on representing analytical decisions explicitly and using LLMs to translate those decisions into code.

3 Methodology

3.1 A schema for describing data analysis

We represent a data analysis, $\mathcal{S} = \{A_1, A_2, \cdots, A_n\}$, as a sequence of n analytical step, A_i, applied to a dataset $\mathcal{D}$, where each step A_i is defined as a tuple:

$$A_i = \langle f_i, \pi_i, r_i, \mathcal{I}_i, \mathcal{O}_i, \omega_i \rangle. \tag{1}$$

The components of A_i are partitioned into a required triplet (f_i, π_i, r_i) and a set of optional components $(\mathcal{I}_i, \mathcal{O}_i, \omega_i)$. The first component of the triplet, $f_i \in \mathcal{F}$, defines the **objective** of a data analysis step that invites multiple justifiable approaches. The second component, $\pi_i \in \Pi$, represents the data analysis **decision** made to achieve the objective f_i, which can be read from both the manuscript and from the code script. The term r_i is the **justification** provided for choosing the decision π_i. This field is usually provided in the manuscript but

not the code script to show the scientific intent of the step. As an example, consider the three steps involved in calculating the Human Development Index (HDI):

1) Scale all four variables (life expectancy, expected years of schooling, mean years of schooling, and GNI per capita) to a [0, 1] range using their corresponding upper and lower limits.
2) Combine the two education-related variables (expected years of schooling and mean years of schooling) into a single education variable.
3) Combine the three core dimensions (health, education, and income) into a single index using a geometric mean.

Structuring this information into a schema is involves defining the objective, decision, and justification for each step. While the decision (what is to be done) and the justification (why it is done) are generally clear, the objective requires identifying what the step is trying to achieve. In other words, it represents the broader goal that could potentially be met through alternative methods (formally introduced in Section 3.2). For instance, in step 3, the objective is to combine the three dimensions. While a simple arithmetic mean could be an alternative solution for this goal, the geometric mean is specifically chosen because it penalizes uneven development among the three dimensions—this reasoning forms the justification. This process results in the basic schema shown in Listing 1:

On the optional components of a data analysis step A_i, the operational flow of the analysis is governed by the input and output spaces, $\mathcal{I}_i$ and $\mathcal{O}_i$, once a dataset is provided. While the output of a step can be a pseudo object that not currently exist, the inputs for any given step must be within the existing data scope, comprising variables in the original dataset and all outputs from previous steps. Formally, we require:

Listing 1 An example schema that specifies the steps in calculating Human Development Index (HDI).

```
meta:
  type: schema
nodes:
- tag: step-scaling
  objective: how to handle variables in different scales?
  decision: apply min-max scaling to each variable
  rationale: to convert the variables to the same scale for future construction of a si
- tag: step-education
  objective: how to combine the two education-related variables?
  decision: take the average
  rationale: the arithmetic mean allows perfect substitutability between the two variab
- tag: step-combine
  objective: how to combine the three dimensions into a single index?
  decision: use the geometric mean
  rationale: the geometric mean penalizes uneven development among the three core dimens
```

$$\mathcal{I}_i \subseteq \left(\mathcal{D} \cup \bigcup_{j<i} \mathcal{O}_j \right), \quad \forall i \in \{2, \ldots, n\}.$$

A schema can also be automatically extracted by Large Language Models (LLMs) from text. If a dataset is provided, the inputs and outputs will also automatically inferred during the extraction. We also introduce a categorical confidence score, ω_i (low, medium, high) that can be useful in case when the source text is ambiguous to guide the LLM to ask for clarification to consolidate the schema extraction.

3.2 Describing alternatives

While one analysis represents a realized series of choices, alternatives can often be taken for to address the same objective, each with its own set of assumptions and implications. For each step i in the sequence, to fulfill each fork f_i, there could be k_i decision choices to choose from: $\{A_i^1, A_i^2, \dots, A_i^{k_i}\}$. We can define a selection vector $r = (r_1, r_2, \dots, r_n)$, where $r_i \in \{1, \dots, k_i\}$, to describe an analysis, comprised of specific series of decisions: $\mathcal{S}^r = \{A_1^{r_1}, A_2^{r_2}, \dots, A_n^{r_n}\}$. A multiverse, used to describe the entire decision spae, is defined as the Cartesian product of these sets: $\mathcal{M} = \{\mathcal{S}^r \mid r_i \in 1, \dots, k_i, \forall i\}$. The cardinality of this space, $|\mathcal{M}| = \prod_{i=1}^{n} k_i$, quantifies the total number of unique, fully-specified analysis.

Returning to the HDI example, in step 3, instead of a geometric mean, an arithmetic mean to combine the three dimensions into a single index. It may not be the best decisions, but still a plausible one to achieve the objective of this step. Listing 2 demonstrates how this single alternative is structured in an alternative YAML file. Additional alternatives can be specified by adding items to the `alternatives` list.

Listing 2 An example YAML file for specifying alternative for Step 3 in calculating Human Development Index (HDI).

```
meta:
  type: alternative
  step: step-combine
  objective: how to combine the three dimensions into a single index?
alternatives:
  - id: arith-mean
    decision: use an arithmetic mean
    rationale: a simple average of the three dimensions
```

3.3 Workflow

The overall workflow for schema-based data analysis is illustrated in Figure 1. A schema can be created either manually from a template or automatically generated by an LLM from a natural-language manuscript. Once created, the schema serves as a structured specification that guides a second LLM in generating the code for the analysis. To improve reliability, our design include a third LLM agent responsible for code validation. While recent LLMs can generate reasonable code based on prompt, they do not consistently achieve perfect accuracy. The validation process runs the generated script in a sandbox environment and if an error occurs, the third LLM agent is invoked to diagnose the script. The validation agent receives the generated code together with the error message adn teh correponding line number produced during execution, and returns a fixed version of the script. The revised script is then re-executed in the sandbox. This cycle continues until the code runs successfully, a maximum number of attempt is reached, or the user terminates the process. As will be show in Section 4, this validation process resolves the majority of R syntax issues encountered during data wrangling and statistical modelling tasks. On top of supporting a single analytical specification, the workflow also allows users to agument a schema with an alternative YAML file. This expands the original schema into a multiverse of parallel schemas, each representing a different set of decision choices. The same code generation and validation workflow is then applied independently to each schema, producing a corresponding set of validation scripts.

3.4 Usage

From an author's perspective: provide it as part of the supplementary materials - verify the results through LLM.

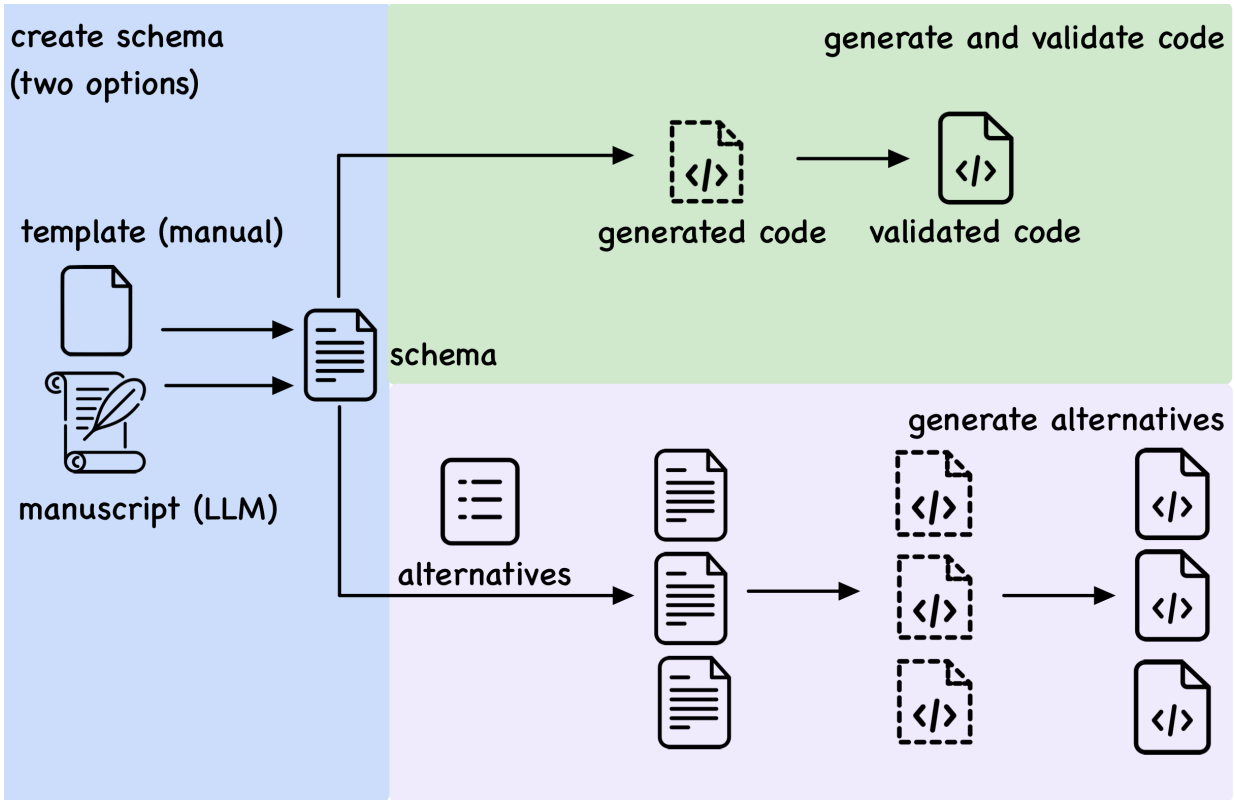


Figure 1: Overview of the schema-driven analysis workflow. Schemas can be authored manually or generated from natural language manuscripts by an LLM. Schemas are used to generate code, which is validated through iterative execution in a sandboxed environment and repaired when errors occur. Alternative specifications can expand a schema into a multiverse of parallel analyses.

from a reader's perspective: having the schema would allow easier exploration of the alternatives, modify code also involves modifying upstream and downstream code, LLM allows it to be automated.

4 Application: many-analysts football study

The many-analysts study posits fixed research questions on a fixed dataset and invites teams of researchers to analyze the problem with whatever software they're comfortable with. In Silberzahn et al. (2018), one question is to answer whether soccer referees are more likely to give red cards to dark-skin-toned player than to light-skin-toned players. The provided dataset contains two raters' rating to the picture of players, the number of redcard received for each player/referee dyade out of the number of games, and a collection of player, referee, and game characteristics: player birthday, height, weight, player position, league country, number of victories, ties, and defeats, and number of yellow and red cards received. Most groups take a modelling approach to estimate the effect of skin rating on the red card issued with varied distribution assumptions (linear, logistic, and poisson) and covariate choices as fixed and random effects. The study is comprised of an initial analysis conducted separately by each team, followed by a round-robin peer evaluation round, and then each group presents their final analysis after incorporating the feedback from the discussion and presents the results in a way they would in a published article with formal Method section. All the materials related to the study is publicly available on the Open Science Framework (OSF).

From the report presented by each team, we can use LLM to convert the decisions into the schema designed and generate the corresponding coding script. From the write-up reports, we extract the sentences relevant to the decisions made in the final analysis by the team.

This step is needed because the original skin rating is 1-5 but later changed to 0-1 and also many teams report their initial analysis and how their final analysis modify on top of the initial analysis. We do this to avoid confusion on reading comprehension with LLM. For general methodology write-up, this step is not necessarily needed [TODO: write it better].

We send the text description of the extractetd decisions for LLM to generate the schema and the code script for each analysis. The code script is run by a separate LLM to validate for potential syntax errors. We manually in-place replace the `glm()` command with `fixest::feglm()` (Bergé et al. 2026) for generalized linear model and the `lme4::glmer()` (Bates et al. 2015) function with `glmmTMB::glmmTMB()` (Brooks et al. 2017) for generalized mix-effect model to speed up the evaluation of eaach script. These two functions are in-place replacements of the classic functions for estimating the relevant model with implementation with multi-threading in C++ (Bergé et al. 2026) and with Template Model Builder (TMB) (McGillycuddy et al. 2025) for better performance. This replacement should not affect the estimation of the effect size. All the LLM-related simulations are conducted with claude-opus-4.5.

In the validation phase, all team scripts were successfully executed. Notably, 10% of the scripts (n = 2) yielded successful results without requiring any automated syntax corrections by the validator, whereas the vast majority (90%) reached completion after two or three iterations. The most frequent correction involved resolving broken and overflow syntax in comments.

Figure 2 shows the agreement of the estimated effect size, in odd ratio, from our experiments (simulated) and the published effect size in Silberzahn et al. (2018) (published). For most teams, the simulation successfully recovers the odd ratio reported by each analyst team, larger variations are seen in the following teams: Team 1, 9, 15, 21, 23, and 27. Here we

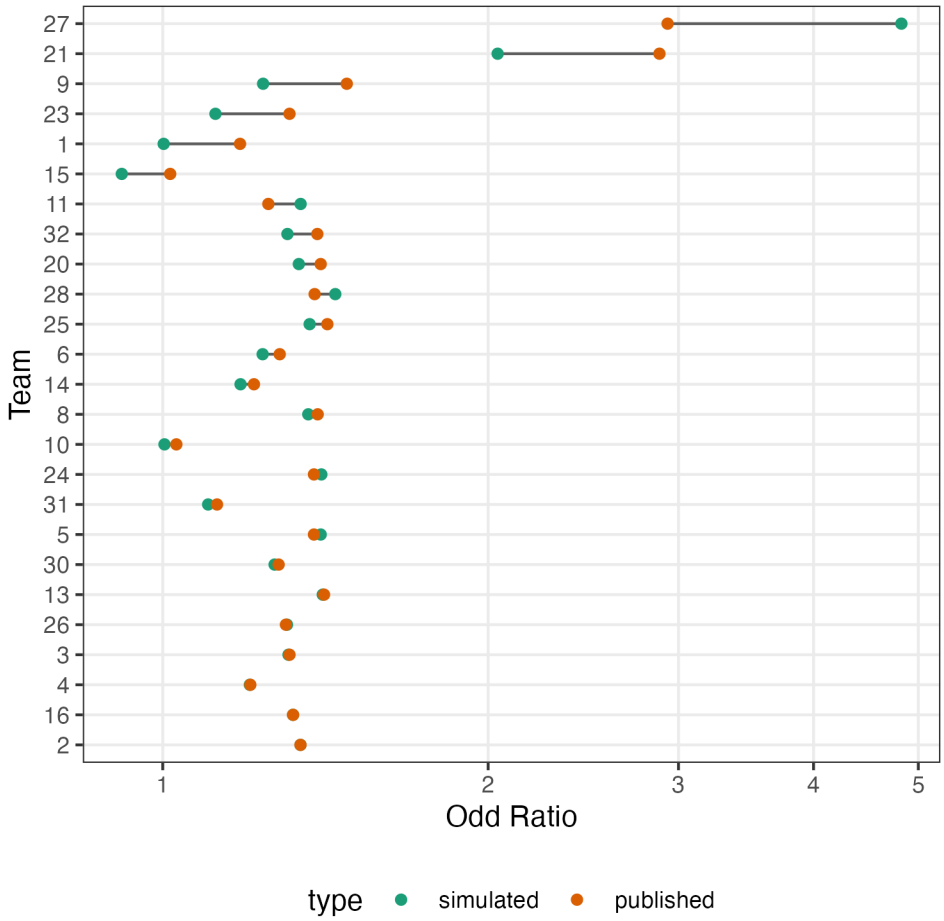


Figure 2: Comparison of the odd ratio of the effect skin color on receiving red cards in football games from our simulation with the schema (simulated) vs. the reported estimated effect size in Silberzahn et al. (2018) (published). Teams are arranged by the difference in odd ratio of our simulation and the reported value. For most teams, the simulation successfully recovers the odd ratio reported by each analyst team, larger variations are seen in the following teams: Team 1, 9, 15, 21, 23, and 27.

listed the reasons for these discrepencies upon investigation:

- **Team 27**: the diffference comes from the ambiguity of the method description, the final approach sections says *In addition to the independent variables specified in the initial approach (skin-tone rating, goals scored, games played, yellow cards, and mean explicit and implicit bias) a variable for player position was included, as well as a variable for the league of the player.* The author is intended to include the interaction effect of rating and position/ league country as other variables in the initial approach, but the LLM interprets it as only include the main effects. This error can be easily caught by comparing the reproducible script provided by the author.
- **Team 9**: the authors mistake player country for referee country in the text description, but their parameter estimation table suggests they are using player country (countryCode - Germany: France) instead.
- **Team 23 and 1**: the discrepency comes from the LLM fails to properly reshape the data from a player-referee dyad per observation to per game observation. The schema properly identifies this step but the generated code doesn't execute this step properly and create a silent error with the validator.
- **Team 15**: the discrpency comes from the default setting for estimating factor variable in R (our similuation) and SPSS (original script). In SPSS, the factor variables are estimated with effect coding, where the coefficients are compared to the grand mean, rather than to the reference group (dummy coding) in R. After changing the `contrast` options, we obtain the reported odd ratio.

4.1 Exploring alternatives

To demonstrate the functionality for specifying alternatives in Section 3.2, we use Group 5 from the football study above as an example. Group 5 uses a binomial generalized linear mixed model (GLMM) to estimate the effect of skin tone rating on the number of red cards received using a logit link function:

$$\textbf{Model 1 (gm1):} \quad \text{logit}(p_{ijk}) = \beta_0 + \beta_1 X_{ijk} + u_j + v_k$$
$$u_j \sim N(0, \sigma^2_{\text{player}}), \quad v_k \sim N(0, \sigma^2_{\text{ref}})$$

$$\textbf{Model 2 (gm2):} \quad \text{logit}(p_{ijk}) = \beta_0 + \beta_1 X_{ijk} + u_j + v_{0k} + v_{1k} X_{ijk}$$
$$u_j \sim N(0, \sigma^2_{\text{player}}), \quad \mathbf{v}_k \sim N\left(\mathbf{0}, \ _{\text{ref}}\right)$$

$$\textbf{Model 3 (gm3):} \quad \text{logit}(p_{ijkm}) = \beta_0 + \beta_1 X_{ijkm} + u_j + v_k + w_{0m} + w_{1m} X_{ijkm}$$
$$u_j \sim N(0, \sigma^2_{\text{player}}), \quad v_k \sim N(0, \sigma^2_{\text{ref}}), \quad \mathbf{w}_m \sim N\left(\mathbf{0}, \ _{\text{country}}\right)$$

where p denotes the probability of a player receiving a red card, and X represents the player's skin tone rating. The subscripts i, j, k, and m index the individual observation, player, referee, and the referee's home country, respectively. The parameter β_0 represents the global baseline log-odds of receiving a red card, and β_1 estimates the fixed effect of a player's skin tone. The random structures account for multi-level variation across groupings with varying intercepts and/or slopes: (1) scalar random intercepts (σ^2) capture baseline shifts for players across all specifications ($u_j \sim N(0, \sigma^2_{\text{player}})$) and for individual referees in Models 1 and 3 ($v_k \sim N(0, \sigma^2_{\text{ref}})$); and (2) variance-covariance matrices ( ) accommodate random slopes, allowing the effect of skin tone to vary across individual referees in Model

2 ($\mathbf{v}_k \sim N(\mathbf{0}, \ _{\text{ref}})$) and across referee countries in Model 3 ($\mathbf{w}_m \sim N\left(\mathbf{0}, \ _{\text{country}}\right)$).

The full workflow of experimenting with alternatives is summarized in **?@fig-alt-workflow**. Alternatives to a specific step are specified in a YAML file in Listing 3, which is then processed to generate multiple copies of the full schema. The code for each schema are then generated separately, and validated to ensure its syntactic correctness and executability. Table 1 summarizes the estimated random effect from the workflow. [TODO: expand more here]

Listing 3 Example of a YAML file specifying alternatives for the random effects structure in the many-analyst analysis of Group 5.

```
meta:
  type: alternatives
  step: specify_random_effects_structure
  fork: how to account for the non-nested cross-classified data structure
alternatives:
  - id: "gm1"
    path: "a random effect of player and referee"
    rationale: ""
  - id: "gm2"
    path: "a random effect of player and referee, and a random effect of skin-tone acros
    rationale: ""
  - id: "gm3"
    path: "a random effect of player, referee, and referees' countries of origin, and a
    rationale: ""
```

Table 1: Summary of random effect estimates for Group 5. The three models (gm1, gm2, and gm3) represent different specifications of the random effects structure to account for the non-nested cross-classified data structure.

group	term	gm1	gm2	gm3
playerShort	var___(Intercept)	0.3737	0.3738	0.2899
refNum	var___(Intercept)	0.1528	0.1932	0.0748
refNum	var___avgrate01	NA	0.0286	NA
refCountry	var___(Intercept)	NA	NA	0.3990
refCountry	var___avgrate01	NA	NA	0.0870

5 Discussion

Our results highlight a persistent gap between computational reproducibility and analytical understandability in modern data analysis workflows. While current reproducibility practices ensure that published results can be regenerated from shared data and code, as analyses grow in complexity and involve iterative decision-making processes, they are less of an idea for readers to understand the analysis. Our proposed schema addresses this gap by introducing an intermediate representation that explicitly encodes analytical decisions to bridge between narrative descriptions and implementation details. When coupled with large language models, this representation enables the extraction of analytical structure from written reports and its translation into analysis scripts, while also supporting the exploration of alternative analytical choices.

To further examine how this workflow behaves in practice, we apply it to the many-analysts

football study. This analysis allows us to identify errors in the original reports and to characterize failure modes of current LLMs to generate code for data science. In particular, we find that LLMs requires a fairly precise specification in order to generate the code correctly for the intended random effect structure. For an educated statistician, the phrase "a random effect of referee number across average rating" would typically imply both a random intercept and a random slope. However, we observe that LLMs interpret this speicification literally as a random slope only. This suggests the gap between the latest models and human analysts in specialized statistical knowledge.

Across the validation process, the LLM successfully corrects a range of implementation issues, including errors caused by overflowing comments and other syntax issues. However, we also observe challenges in data reshaping tasks when the transformation requires changes in the unit of analysis. In the study, some teams decide to transform the data from the original aggregated player-referee dyad-level format to game-level format. This transformation requires disaggregating aggregated outcome counts (win, loss, and draw) into individual game records, changing the unit of analysis from dyads to matches. We also observe some difficulties in parsing birth dates. Since the LLM does not have access to the underlying raw data, due to privacy constraint, it must rely on trial and error when inferring the correct date formats ("DD-MM-YYYY."). This suggests that certain types of data analysis tasks may not be fully automatable in settings where the model cannot directly inspect the data, even if the data is theoretically available for computation.

TODO: how do we modularize data analysis

Appendix

Built-in prompts

Prompt: Extracting the schema from the text

```
You are an expert methodologist and data pipeline architect. I will
provide a text describing a multiverse analysis and a summary of the
actual dataset being used.


=== YOUR TASK ===


Extract a chronological list of methodological decisions (nodes) AND
map the exact data flow (inputs/outputs) for each node
simultaneously.


=== RULES ===


1. THEORY EXTRACTION: For each node, extract:


- 'id': A unique snake_case identifier.


- 'fork': MUST be framed as an open methodological goal that invites
multiple possible approaches. It must NOT describe the final choice.


- 'path': A 'path' is strictly a POSITIVE methodological decision
```

```
that has potential theoretical alternatives, chosen to resolve the
'fork'.

- 'rationale': WHY that decision (the path) was made, extracted from
the text.

2. DATA MAPPING: Assign 'inputs' (EXACT column names from the dataset
OR outputs from previous nodes) and 'outputs' (invented snake_case
objects like 'df_clean' or 'ranef_spec').

3. ANTI-ABSTRACTION (CRITICAL): If the text lists specific variables
(e.g., 'centered rater, meanIAT'), DO NOT summarize them away. You
MUST capture those specific variables in the 'inputs' array.

4. INLINE ARRAYS: Format arrays strictly on one line: `inputs: [var1,
var2]`.

5. CONFIDENCE & CLARIFICATION: Rate your mapping confidence (HIGH,
MEDIUM, LOW). If the text abstracts a step and you cannot confidently
match it to specific dataset columns, set confidence to LOW and
autogenerate a 'clarification_question' asking the user which exact
columns to use. If HIGH, output 'null'.

6. Output ONLY valid YAML without markdown formatting.
```

```
=== REQUIRED YAML STRUCTURE EXAMPLE ===

meta: type: schema nodes: - fork: variables are in different scales
path: apply min-max scaling to each variable justification: to put
them on the same scale for combination id: step-scaling confidence:
high - fork: combine the school variables into one dimension path:
average exp sch and avg sch justification: the most intuitive way id:
step-education confidence: low

=== DATASET SUMMARY ===

NA

=== METHODOLOGY TEXT ===

[inset text here]
```

Prompt: generating code from the schema

```
You are an expert R programmer. Attached is a text document
containing a SCHEMA that defines a data processing pipeline.

No data file is specified - if data loading is required, use a
sensible placeholder or note it in a comment. Your task is to write a
complete, working R script that implements this pipeline
```

step-by-step. Use modern R practices (like dplyr or base pipe) and ensure variables flow correctly from one step to the next as defined by the inputs and outputs. Output ONLY the complete R script. Do not start with markdown formatting blocks (like ```R) or backticks.

Prompt: validating the code generated

You are an expert R developer and automated debugging agent. The user has provided a data file located at the relative path: `data/[insert the path to the data]'. Your purpose is to receive broken R scripts, diagnose the execution error, and return the corrected script.

CRITICAL INSTRUCTIONS: 1. Fix the bug while preserving the original intent and logic of the script. 2. If the error is a missing package, add the necessary `library()` call at the top. 3. If the error is a missing variable, ensure it is properly initialized before use. 4. Do NOT hallucinate new data files or external dependencies unless explicitly provided in the context.

STRICT OUTPUT CONSTRAINTS: - You must output ONLY valid, fully executable R code. - Absolutely NO markdown formatting. Do NOT wrap your response in ```R or ```. - Absolutely NO conversational filler, greetings, explanations, or comments about what you fixed. - The first character of your response must be R code, and the last character must be R code.

Prompt: generating alternatives for a specific step

```
You are an expert Data Analyst and Methodologist. You are reviewing
an analysis schema to identify "Forking Paths" -- alternative
analytical choices that are equally valid but distinct from the
current approach.


=== DEFINITIONS ===


The schema provided to you consists of steps with:


- **ACTION**: The goal of the step (What needs to be done).


- **DECISION**: The specific implementation chosen (How it is done).


- **JUSTIFICATION**: The reasoning behind that decision.


- **ID**: The unique identifier for the step (kebab-case).


=== TASK ===


Focus specifically on the step tagged: "[inset step number here]".
Your goal is to generate 3 distinct, valid alternatives for this
step.
```

For each alternative: 1. **Keep the same ACTION** (the goal remains constant).

2. **Change the DECISION** to a different but methodologically sound approach.

3. **Provide a new JUSTIFICATION** explaining why this alternative is valid.

4. **Create a new ID** that reflects the new decision (must be kebab-case).

=== OUTPUT FORMAT ===

Please output the result in **strictly valid YML format**.

Crucial Formatting Rules:

1. Include a `meta` section at the top with `type: alternative` and the `step`.

2. Output strictly valid YML. All text values (decision, justification) must be enclosed in double quotes ("). Do not use block styles (| or >). Do not wrap lines or insert \n characters

```
within the quotes; keep the text as a single continuous string.


3. Do not include markdown code fences (like ```yml) or
conversational text. Just the raw YML.


=== REQUIRED YML STRUCTURE EXAMPLE ===


meta: type: tines_alternative step: [inset step number here]
alternatives: - id: step-new-method-name action: Repeat the original
action decision: "Description of the new decision..."  justification:
"This is the reasoning for why this alternative is valid."  - id:
step-another-method ...
```

References

Allaire, J., Teague, C., Scheidegger, C., Xie, Y., Dervieux, C. & Woodhull, G. (2026), 'Quarto'.
URL: *https://github.com/quarto-dev/quarto-cli*

Bates, D., Mächler, M., Bolker, B. & Walker, S. (2015), 'Fitting linear mixed-effects models using lme4', *Journal of Statistical Software* **67**(1), 1–48.

Bergé, L. R., Butts, K. & McDermott, G. (2026), 'Fast and user-friendly econometrics estimations: The R package fixest', *arXiv preprint arXiv:2601.21749* .

Bertran, M., Fogliato, R. & Wu, Z. S. (2026), 'Many AI Analysts, One Dataset: Navigating

the Agentic Data Science Multiverse'.

URL: *http://arxiv.org/abs/2602.18710*

Brooks, M. E., Kristensen, K., van Benthem, K. J., Magnusson, A., Berg, C. W., Nielsen, A., Skaug, H. J., Maechler, M. & Bolker, B. M. (2017), 'glmmTMB balances speed and flexibility among packages for zero-inflated generalized linear mixed modeling', *The R Journal* **9**(2), 378–400.

Chen, M., Tworek, J., Jun, H. et al. (2021), 'Evaluating large language models trained on code', *arXiv preprint arXiv:2107.03374* .

URL: *https://arxiv.org/abs/2107.03374*

Gao, R. & Xiao, S. C. (2026), 'Nonstandard Errors in AI Agents'. arXiv:2603.16744 [cs].

URL: *http://arxiv.org/abs/2603.16744*

Hu, S., Lu, C. & Clune, J. (2024), Automated Design of Agentic Systems.

URL: *https://openreview.net/forum?id=t9U3LW7JVX*

Huang, J., Wang, C., Zhang, J., Yan, C., Cui, H., Inala, J. P., Clement, C. & Duan, N. (2022), Execution-based Evaluation for Data Science Code Generation Models, *in* E. Dragut, Y. Li, L. Popa, S. Vucetic & S. Srivastava, eds, 'Proceedings of the Fourth Workshop on Data Science with Human-in-the-Loop (Language Advances)', Association for Computational Linguistics, Abu Dhabi, United Arab Emirates (Hybrid), pp. 28–36.

URL: *https://aclanthology.org/2022.dash-1.5/*

Jansen, J. A., Manukyan, A., Khoury, N. A. & Akalin, A. (2025), 'Leveraging large language models for data analysis automation', *PLOS ONE* **20**(2), e0317084.

URL: *https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0317084*

Jiang, J., Wang, F., Shen, J., Kim, S. & Kim, S. (2026), 'A survey on large language mod-

els for code generation', *ACM Transactions on Software Engineering and Methodology* **35**(2), 1–72.

Jun, E., Birchfield, M., De Moura, N., Heer, J. & Just, R. (2022), 'Hypothesis formalization: Empirical findings, software limitations, and design implications', *ACM Transactions on Computer-Human Interaction (TOCHI)* **29**(1), 1–28.

Jun, E., Seo, A., Heer, J. & Just, R. (2022), Tisane: Authoring statistical models via formal reasoning from conceptual and data relationships, CHI '22, Association for Computing Machinery, New York, NY, USA, p. 1–16.
URL: *https://dl.acm.org/doi/10.1145/3491102.3501888*

Kluyver, T., Ragan-Kelley, B., Pérez, F., Granger, B., Bussonnier, M., Frederic, J., Kelley, K., Hamrick, J., Grout, J., Corlay, S., Ivanov, P., Avila, D., Abdalla, S., Willing, C. & development team, J. (2016), Jupyter notebooks - a publishing format for reproducible computational workflows, *in* F. Loizides & B. Scmidt, eds, 'Positioning and Power in Academic Publishing: Players, Agents and Agendas', IOS Press, Netherlands, pp. 87–90.
URL: *https://eprints.soton.ac.uk/403913/*

Kohler, B., Zollikofer, D., Einsiedler, J., Hoyle, A. & Ash, E. (2026), 'Read the Paper, Write the Code: Agentic Reproduction of Social-Science Results'.
URL: *http://arxiv.org/abs/2604.21965*

Lai, Y., Li, C., Wang, Y., Zhang, T., Zhong, R., Zettlemoyer, L., Yih, W.-T., Fried, D., Wang, S. & Yu, T. (2023), DS-1000: A Natural and Reliable Benchmark for Data Science Code Generation, *in* 'Proceedings of the 40th International Conference on Machine Learning', PMLR, pp. 18319–18345.

URL: *https://proceedings.mlr.press/v202/lai23b.html*

Leisch, F. (2002), Sweave: Dynamic generation of statistical reports using literate data analysis, *in* W. Härdle & B. Rönz, eds, 'Compstat 2002 — Proceedings in Computational Statistics', Physica-Verlag GmbH, Heidelberg, pp. 575–580.

Luo, A., Du, J., Tian, F., Xian, X., Specht, R., Wang, G., Bi, X., Fleming, C., Srinivasa, J., Kundu, A., Hong, M. & Ding, J. (2025), Can Agentic AI Match the Performance of Human Data Scientists?, *in* '2025 IEEE 10th International Workshop on Computational Advances in Multi-Sensor Adaptive Processing (CAMSAP)', pp. 206–210. ISSN: 2994-8983.

URL: *https://ieeexplore.ieee.org/document/11423949/*

McGillycuddy, M., Warton, D. I., Popovic, G. & Bolker, B. M. (2025), 'Parsimoniously fitting large multivariate random effects in glmmTMB', *Journal of Statistical Software* **112**(1), 1–19.

Owusu Boakye, K. D. N. (2026*a*), *chatLLM: A Flexible Interface for 'LLM' API Interactions*. R package version 0.1.4.

URL: *https://CRAN.R-project.org/package=chatLLM*

Owusu Boakye, K. D. N. (2026*b*), *LLMAgentR: Language Model Agents in R for AI Workflows and Research*. R package version 0.3.2.

URL: *https://CRAN.R-project.org/package=LLMAgentR*

Peng, R. D. (2011), 'Reproducible Research in Computational Science', *Science* **334**(6060), 1226–1227.

URL: *https://www.science.org/doi/10.1126/science.1213847*

Shah, S. M. H., Hopfgartner, F. & Bleier, A. (2026), 'Automating Computational Repro-

ducibility in Social Science: Comparing Prompt-Based and Agent-Based Approaches'. arXiv:2602.08561 [cs].

URL: *http://arxiv.org/abs/2602.08561*

Silberzahn, R., Uhlmann, E. L., Martin, D. P., Anselmi, P., Aust, F., Awtrey, E., Bahník, Š., Bai, F., Bannard, C., Bonnier, E., Carlsson, R., Cheung, F., Christensen, G., Clay, R., Craig, M. A., Dalla Rosa, A., Dam, L., Evans, M. H., Flores Cervantes, I., Fong, N., Gamez-Djokic, M., Glenz, A., Gordon-McKeon, S., Heaton, T. J., Hederos, K., Heene, M., Hofelich Mohr, A. J., Högden, F., Hui, K., Johannesson, M., Kalodimos, J., Kaszubowski, E., Kennedy, D. M., Lei, R., Lindsay, T. A., Liverani, S., Madan, C. R., Molden, D., Molleman, E., Morey, R. D., Mulder, L. B., Nijstad, B. R., Pope, N. G., Pope, B., Prenoveau, J. M., Rink, F., Robusto, E., Roderique, H., Sandberg, A., Schlüter, E., Schönbrodt, F. D., Sherman, M. F., Sommer, S. A., Sotak, K., Spain, S., Spörlein, C., Stafford, T., Stefanutti, L., Tauber, S., Ullrich, J., Vianello, M., Wagenmakers, E.-J., Witkowiak, M., Yoon, S. & Nosek, B. A. (2018), 'Many analysts, one data set: Making transparent how variations in analytic choices affect results', *Advances in Methods and Practices in Psychological Science* **1**(3), 337–356. Publisher: SAGE Publications Inc.

URL: *https://doi.org/10.1177/2515245917747646*

Wickham, H., Cheng, J., Jacobs, A., Aden-Buie, G. & Schloerke, B. (2025), *ellmer: Chat with Large Language Models*. R package version 0.4.0.

URL: *https://CRAN.R-project.org/package=ellmer*

Wrobel, J., Hector, E. C., Crawford, L., McGowan, L. D., da Silva, N., Goldsmith, J., Hicks, S., Kane, M., Lee, Y., Mayrink, V., Paciorek, C. J., Usher, T. & Wolfson, J. (2024), 'Partnering With Authors to Enhance Reproducibility at JASA', *Journal of the American Statistical Association* **119**(546), 795–797. _eprint:

https://doi.org/10.1080/01621459.2024.2340557.

URL: *https://doi.org/10.1080/01621459.2024.2340557*

Xie, Y. (2015), *Dynamic Documents with R and knitr*, 2nd edn, Chapman and Hall/CRC, Boca Raton, Florida. ISBN 978-1498716963.

URL: *https://yihui.org/knitr/*

Xu, Y. & Yang, L. Y. (2026), 'Scaling Reproducibility: An AI-Assisted Workflow for Large-Scale Replication and Reanalysis'.

URL: *http://arxiv.org/abs/2602.16733*

Zhang, H. S. & Peng, R. D. (2025), 'Inside Out: Externalizing Assumptions in Data Analysis as Validation Checks', *Journal of Data Science* pp. 1–18.